

Disentangling Confounded Groups

Anonymous Authors¹

Abstract

Group-instance disentanglement is the problem of learning separate representations for within-group and across-group variation. In our work, we propose the Context-Aware Variational Autoencoder (CxVAE), a method that can perform group-instance disentanglement on data with “emergent” groups (i.e. data where the group and instance variables cannot be inferred from a single datapoint). First, we introduce a simple synthetic dataset to show that current state-of-the-art methods based on the Group VAE fail to learn disentangled representations on datasets with emergent groups. Next, we present a modified model that 1) infers the instance variable conditionally on the group variable, 2) has a more expressive group encoder, 3) optimizes a regularization objective that encourages independence between the instance variable and the grouping. Finally, we show empirically the impact of each modification with respect to holdout data likelihood, unsupervised translation, and other disentanglement metrics.

1. Introduction

Group-instance disentanglement is the task of learning a representation network $r(\mathbf{x})$ which transforms a group of observations $\mathbf{x} = \{x_1, \dots, x_K\}$ (e.g. images grouped by author, temperature readings grouped by weather station, etc) into a *group* code u and a set of *instance* codes $\mathbf{z} = \{z_1, \dots, z_K\}$. The goal is that the group code should capture all the information that distinguishes one group from another, while the instance code should capture all the information that distinguishes the observations within a group, but none of the group information.

All the current state-of-the-art approaches to group-instance disentanglement (Bouchacourt et al., 2018; Hosoya, 2019; Shu et al., 2020; Chen & Batmanghelich, 2020; Locatello




Figure 1. Probabilistic graphical model of the Group Variational Autoencoder. u are group variables and v are instance variables. The conditional $q(v_{ik}|u_i)$ (shown in blue) is absent from the existing literature but present in our proposed model. This enables our model to accurately infer the instance variable even when the groups are confounded.

et al., 2020) follow the Variational Autoencoder framework (Kingma & Welling, 2014; Rezende et al., 2014) whereby a generative model mapping latent variables (group and instance) to data observations is trained using variational inference. We call this family of models the Group Variational Autoencoder (GVAE). These methods vary in how they define the variational latent posterior and any additional regularization objectives they might have.

We believe that GVAE has the potential to advance many scientific fields where analysing the causal factors of grouped data is essential, including medicine (Liang & Zeger, 1986), biology (Quinn & Keough, 2002) and psychology (Schafer & Graham, 2002). That analysis is typically performed using mixed-effects models (Pinheiro & Bates, 2001), which decompose the data variation into random-effects (group variables) and individual error terms (instance variables). However, these parameters can only be estimated efficiently in very simple models. By contrast, the GVAE offers a highly expressive mixed-effects model (the decoder) and an inbuilt inference network (the encoder).

In order to learn useful disentangled representations for

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

these problems, the GVAE has to capture the joint distribution between the group and instance variables conditioned on the data. In other words, the representation network must be able to distinguish between two identical observations coming from different groups, thus encoding them with different instance variables. We call data groups of this kind “emergent”, because the group and instance variables are emergent properties of the group as a whole, rather than any individual observation.

Emergent groups are very common, occurring in datasets where the latent variables cannot be accurately inferred from a single observation. For example, in the context of soybean plant growth (Davidian & Giltinan, 2003), it is impossible to deduce the age (instance) of a plant based solely on its leaf weight (observation) without also knowing its genotype (group).

Unfortunately, the current state-of-the-art methods fail to learn disentangled representations on emergent groups. We show this empirically by creating a simple dataset of student exam scores grouped by school. The groups in this dataset are emergent, because the instance variable captures the score of a student in relation to the distribution of scores at their school. Instead of disentangling, the existing methods mostly forego the group variable and use the instance variable to encode the variation across groups. The main

In order to tackle this challenge, we propose a new GVAE model with several novel modifications. The main novelty is an instance encoder conditional on the previously inferred group variable, which enables the variational latent posterior to match the generative latent posterior. Additionally, we use an expressive Deep Set network (Zaheer et al., 2017) to infer the group variable more accurately, and also design a novel regularization objective that encourages independence between the inferred instance variable and the data group.

Our work comprises the following contributions:

1. We formalize the notion of “emergent” groups and create a simple synthetic dataset on which the current GVAE methods fail to learn disentangled representations.
2. We propose a new model, called the Context-Aware Variational Autoencoder (CxVAE) which can learn disentangled representations on emergent groups.
3. We perform extensive evaluation showing that our model outperforms the existing methods in holdout-data log-likelihood, unsupervised translation, and ground-truth factor prediction.
4. We perform an ablation study showing the relative performance gain of each modification to the GVAE.

2. Generative Group-Instance

The Group Variational Autoencoder (Bouchacourt et al., 2018; Hosoya, 2019) is a family of models that use two latent variables to represent grouped data: one that captures the variation within groups, and one for the variation across groups.

Assume a dataset of the form $\{x_{ik}\}_{i \in 1:N, k \in 1:K_i}$ where N is the number of groups and K_i is the number of observations in group i . GVAE defines a generative model that maps a $\mathcal{N}(0, 1)$ group latent variable u_i and a $\mathcal{N}(0, 1)$ instance latent variable v_{ik} to a given data observation x_{ik} . In other words, the likelihood of a group is:

$$p(\mathbf{x}) = \mathbb{E}_{p(u)} \prod_{k=1}^K \mathbb{E}_{p(v_k)} [p(x_k|u, v_k)]$$

We omit the index of the group i for notational simplicity, since the groups are independent and identically distributed.

2.1. Variational Inference

Because the exact likelihood is intractable, the Variational Autoencoder (Kingma & Welling, 2014; Rezende et al., 2014) performs optimization by introducing a variational latent posterior $q(u, \mathbf{v}|\mathbf{x})$ and maximizing the Evidence Lower Bound (Jordan et al., 2004):

$$\log p(\mathbf{x}) \geq \mathbb{E}_{q(u, \mathbf{v}|\mathbf{x})} [\log p(\mathbf{x}|u, \mathbf{v})] - \text{KL}[q(u, \mathbf{v}|\mathbf{x})||p(u, \mathbf{v})]$$

The models in the GVAE family use a class of variational distributions that assume independence between the latent variables in a group when conditioned on the data.

$$q(u, \mathbf{v}|\mathbf{x}) = q(u|\mathbf{x}) \prod_{k=1}^K q(v_k|x_k)$$

In our work, we show that this assumption hinders disentanglement when the generative model is entangled.

2.2. Group Encoder

The variational group posterior is realised as normal density with μ, Σ computed with an encoder network. The way to implement the encoder network is not obvious, since the number of inputs K varies across groups. Hosoya (2019) encode separately each observation in the group using the same encoder E_u to produce μ_k, Σ_k and then averages the outputs.

$$\mu = \frac{1}{N} \sum_{k=1}^K \mu_k, \Sigma = \frac{1}{N} \sum_{k=1}^K \Sigma_k$$

Bouchacourt et al. (2018) also encode each observation individually and then accumulate the evidence through a product of normal densities, computed using the following equations:

$$\Sigma^{-1} = \sum_{k=1}^K \Sigma_k^{-1}, \quad \mu^T \Sigma^{-1} = \sum_{k=1}^K \mu_k^T \Sigma_k^{-1}$$

They justify that such a product of normals produces a valid evidence accumulation using the following result:

$$q(u|\mathbf{x}) \propto \prod_{k=1}^K q(u|x_k)$$

However, the above is not a universal property, since

$$\begin{aligned} q(u|\mathbf{x}) &= \frac{\prod_{k=1}^K q(x_k)}{q(\mathbf{x})} \frac{1}{q(u)^{K-1}} \prod_{k=1}^K q(u|x_k) \\ &\propto \frac{1}{q(u)^{K-1}} \prod_{k=1}^K q(u|x_k) \end{aligned} \quad (1)$$

In fact, by using a product of normals to accumulate evidence, the authors implicitly assume that the marginal distribution of the inferred group variable is a uniform. This has the effect of sampling u values which are less representative of the current group and more skewed towards the marginal distribution of u .

In our work, we propose a more general approach to encoding u by using a Deep Sets network (Zaheer et al., 2017) to encode the whole set of observations instead of encoding each observation separately.

2.3. Regularization

In certain cases, the model might learn to encode both kinds of variation (within- and across-group) in the instance variable, effectively turning the model into a standard VAE. In this eventuality, the group variable becomes irrelevant and disentanglement is not achieved.

Such behaviour has been identified by Hosoya (2019) and Németh (2020) to occur when the instance code too high-dimensional, the instance encoder too expressive, or group sizes too small. One solution is to limit the dimensionality of the instance code (Hosoya, 2019), with the downside of hindering the overall model performance.

As a more targeted solution, Hosoya (2019) propose an adversarial loss minimizing the mutual information between an observation and the instance variable inferred from the other observations in the group:

$$I_r(x, v) = \text{KL}[r(x, v) || r(x)r(v)]$$

where $r(x, v) = r(v|x)r(x)$ is the joint distribution of an observation and the instance variable inferred from any of the other observations in the group and

$$r(v|x_k) = \frac{1}{K-1} \sum_{l=1, l \neq k}^K q(v|x_l)$$

The mutual information is approximated empirically using the results of Belghazi et al. (2018).

$$\begin{aligned} I_r(x, v) &\approx \max_T \mathbb{E}_{r(x, v)} [T(x, v)] \\ &\quad - \log \mathbb{E}_{r(x)r(v)} [\exp T(x, v)] \end{aligned} \quad (2)$$

T is a neural network and the expectation terms are computed by sampling.

- To sample $r(x, v)$, first choose a group i , then choose two instances from that group $k, l \in K_i$. x will be the observation x_{ik} and v will be sampled from $q(v|x_{il})$.
- To sample $r(x)r(v)$, choose two groups i, j and two instances in each group $k \in K_i, l \in K_j$. Take x_{ik} for x and sample $q(v|x_{jl})$.

In our view, this method has the following limitation: Even when the instance variable does contain group information, the value of $I_r(x, v)$ might still be small, because it might be difficult to ascertain the group based on one single observation x . In our work, we propose a modification to this regularization term such that the network T takes as input all the observations in the group instead of only one.

3. Entangled Group and Instance Variables

We call the group and instance variables *entangled* when they are not independent conditioned on the data $p(u, v|x) \neq p(u|x)p(v|x)$. A useful heuristic for establishing whether the variables are entangled is to ask ‘‘Does knowing the group variable for an observation influence my belief about its instance variable?’’

This property of the generative model is present in many machine learning tasks, such as collaborative filtering, 3D novel view synthesis. For example, in the context of the Netflix Challenge, where the task was to predict what score a user would give to a new film, one cannot infer what film is associated with a given score without also knowing the user.

Strictly speaking, most real-world models are entangled. However, in many cases, the mutual information between the group and instance variable, conditioned on the observation, is negligible. For example, in handwritten digit recognition, one can infer the digit value depicted in an image without knowing the author.

In this paper, we claim that the current methods in the GVAE family do not perform well in tasks where the group and instance variables are entangled.

3.1. Exam-Scores Problem

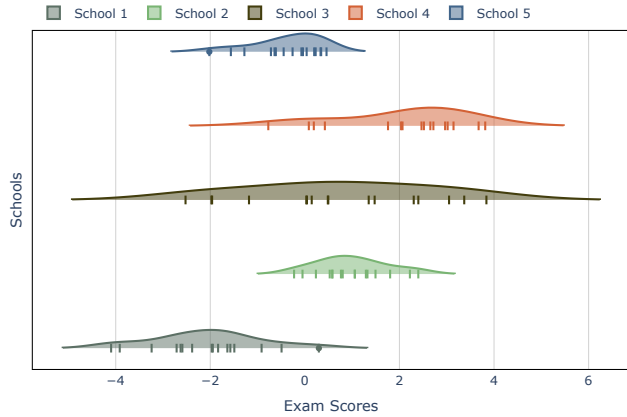


Figure 2. Normalized exam scores of individual students grouped by school.

Suppose we wanted to model the exam scores of students from different schools. Our model must separate the school-level effect (the group factor) from the student-level effect (the instance factor). We define the following generative model:

$$x_{ik} = 2\mu_i + (\sigma_i^2 + 1)v_{ik} + \epsilon_{ik}, \quad i \in 1:N, \quad k \in 1:K_i$$

where x_{ij} is the student score, $u_i = (\mu_i, \sigma_i)$ is the school-level effect, v_{ij} is the student-level effect, and ϵ_{ik} is a normally-distributed error term. We assume a $\mathcal{N}(0, 1)$ prior distribution for the latent variables.

We first sample the model to generate a dataset ($N = 32,768$, $K_i \sim \text{Poisson}(16) + 8$) and then use the same model as the generative model in our Variational Autoencoder, instead of a neural network. The figure below shows what the data looks like.

Looking at the data, it is easy to see that this model is entangled, because the relative performance v of a student within their own school, given their absolute score x , depends on the distribution of scores within the school u .

4. Context-Aware Variational Autoencoder (CxVAE)

We propose a new model which can perform well on datasets generated from entangled group and instance variable. We call our model the Context-Aware Variational Autoencoder. Our model comprises the following changes with respect to the standard GVAE:

1. The group encoder is implemented as a Deep Sets network (Zaheer et al., 2017). The encoder has the following form:

$$\mu, \Sigma = E_u(\mathbf{x}) = E_u^B \left(\sum_{k=1}^K E_u^A(x_k) \right)$$

where E_u^A, E_u^B are two neural networks.

2. The variational instance posterior is dependent on the inferred group variable:

$$q(u, \mathbf{v}|\mathbf{x}) = q(u|\mathbf{x}) \prod_{k=1}^K q(v_k|x_k, u)$$

In practice, our instance encoder takes as input a vector concatenating x_k and u . This idea is not new, and has been used previously in sequence disentanglement (Li & Mandt, 2018). This allows the instance encoder to differentiate the between observations with similar values but which come from different groups.

3. We propose a regularization objective similar to the one in N  meth (2020), but whereby we minimize the mutual information between one inferred instance variable and all the other observations in the group. More precisely, our objective is to minimize $I_r(\mathbf{x}_{-k}, v_k)$ where $r(v_k|\mathbf{x}_{-k}) = q(v_k|x_k, u)$. Following the same approximation as in N  meth (2020), the objective takes the following form:

$$I_r(\mathbf{x}_{-k}, v_k) \approx \max_T \mathbb{E}_{r(\mathbf{x}_{-k}, v_k)} [T(\mathbf{x}_{-k}, v_k)] - \log \mathbb{E}_{r(\mathbf{x}_{-k})r(v_k)} [\exp T(\mathbf{x}_{-k}, v_k)] \quad (3)$$

T is implemented as a Deep Sets neural network for the observations, with the instance code concatenated in the middle:

$$T(\mathbf{x}_{-k}, v_k) = T^\beta \left(v_k, \frac{1}{K} \sum_{l=1, l \neq k}^K T^\alpha(x_l) \right)$$

Again, the expectations are computed by sampling:

- To sample $r(x, v)$, first choose a group i , then choose one instance from that group $k \in K_i$. $\mathbf{x}_{i,-k}$ will be all the observations in the group apart from k , and v_k will be sampled from $q(v|x_{ik})$.

- To sample $r(\mathbf{x}_{-k})r(v_k)$, choose two groups i, j and two instances in each group $k \in K_i, l \in K_j$. Take $\mathbf{x}_{i,-k}$ for \mathbf{x}_{-k} and sample $q(v|v_{jl})$.

This regularization objective also minimizes the mutual information between the inferred instance variables and the true data generating group variable, but uses as a proxy for the latter all but one of the observations in the group, instead of just one observation.

5. Measuring Disentanglement

In the context of the GVAE family, disentanglement is a property of the variational latent posterior. The inferred group and instance variables are disentangled when they are maximally informative about the group and instance variables of the true data generating distribution. We assume this true model has the same factorization of the joint distribution as the generative model, but the parameters are unknown.

Therefore, we can use the mutual information between the true and inferred latents as a measure of disentanglement.

$$\begin{aligned} I(U^t; U^q) &= \text{KL}[\text{Pr}_{U^q, U^t} || \text{Pr}_{U^q} \text{Pr}_{U^t}] \\ &= \mathbb{E}_{U^q} \mathbb{E}_{U^t | U^q} \left[\log \frac{\text{Pr}_{U^t | U^q}}{\text{Pr}_{U^t}} \right] \end{aligned}$$

and since Pr_{U^t} does not depend on q

$$= \mathbb{E}_{U^q} \mathbb{E}_{U^t | U^q} [\log \text{Pr}_{U^t | U^q}] + \text{const}$$

We estimate the density $\text{Pr}_{U^t | U^q}$ by training a regression network to predict the value of U^t given U^q . The mean-squared error of the prediction on the holdout set is a monotonic function of the mutual information with respect to changes in q .

5.1. Unsupervised Translation

Unsupervised translation is the process of transforming an observation by changing its group code while keeping its instance code fixed. This is a common downstream task for disentangled representations because it requires a clean separation between the group and instance representations (Tenenbaum & Freeman, 2000). Therefore, it can be used to quantify the quality of disentanglement.

Formally, let i, j be the indices of the source and target group, respectively. Translation involves sampling a group of instance codes from the source group $q(\mathbf{v}_i | \mathbf{x}_i)$ and a group code from the target group $q(u_j | \mathbf{x}_j)$. Then, we generate the translated observations by combining each instance code with the group code $p(\mathbf{x}'_k | u_j, \mathbf{v}_j)$.

We measure translation quality by taking the mean-squared error between the translation performed with each model

and the ground-truth translation computed using the ground-truth factors and the true data-generating process.

6. Evaluation

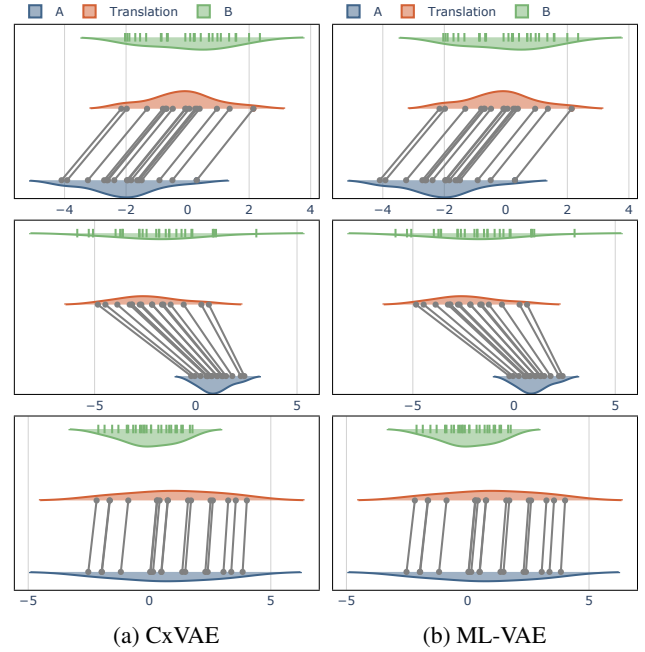


Figure 3. Qualitative comparison of unsupervised translation between CxVAE (top) and ML-VAE (Bouchacourt et al., 2018) (bottom). Our model captures the distribution of values in the target group better than the existing methods.

We test CxVAE against the other models from the GVAE family.

6.1. Quantitative Evaluation

6.2. Ablation Study

In order to quantify the effect of each proposed improvement, we perform an ablation study whereby we measure the decrease in performance resulting from replacing a proposed element of our model with a current alternative.

The proposed improvements have the best individual performance with respect to both reconstruction and translation error in each of the three categories. The choice of a group-aware instance encoder leads to the most significant increase in performance.

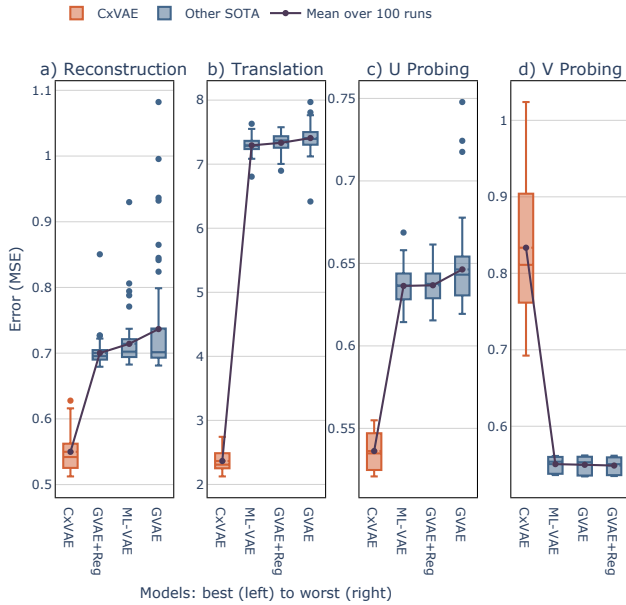


Figure 4. Our CxVAE produces significant improvements over the state-of-the-art in every disentanglement metric considered: hold-out reconstruction error (**lower is better**), translation error (**lower is better**), *U*-probing (**lower is better**) and *V*-probing (**higher is better**). We show box-plots for the distribution of scores over 100 random initialisations.

6.3. Relative strength of the group and instance components

References

- Belghazi, M. I., Baratin, A., Rajeswar, S., Ozair, S., Bengio, Y., Hjelm, R. D., and Courville, A. C. Mutual information neural estimation. In *ICML*, 2018.
- Bouchacourt, D., Tomioka, R., and Nowozin, S. Multi-level variational autoencoder: Learning disentangled representations from grouped observations. *AAAI Conference on Artificial Intelligence*, 2018.
- Chen, J. and Batmanghelich, K. Weakly supervised disentanglement by pairwise similarities. *AAAI Conference on Artificial Intelligence*, 2020.
- Davidian, M. and Giltinan, D. M. Nonlinear models for repeated measurement data: An overview and update. Number 4, pp. 387–419, 2003.
- Hosoya, H. Group-based learning of disentangled representations with generalizability for novel contents. In *IJCAI*, 2019.
- Jordan, M. I., Ghahramani, Z., Jaakkola, T., and Saul, L. K.

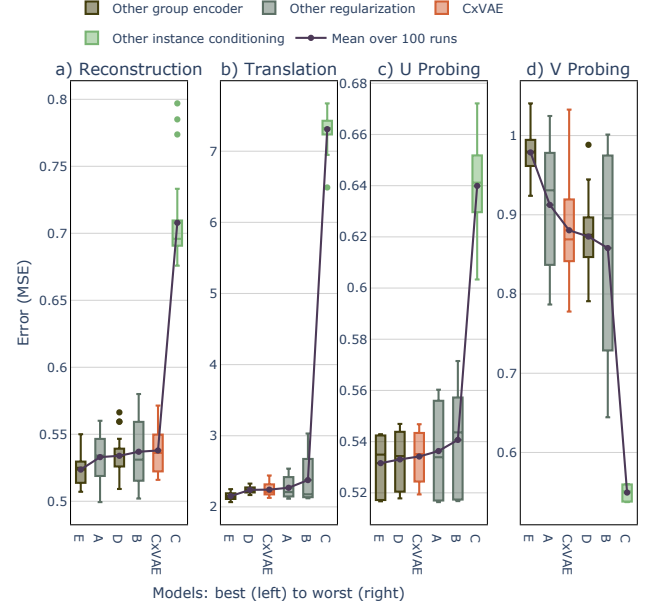


Figure 5. **Ablation study.** Replacing any novel component of our CxVAE with an existing alternative leads to worse performance. Note that the largest reduction in performance is brought by making the instance encoder unconditional. The next is changing the regularization, followed by using a less expressive group encoder.

An introduction to variational methods for graphical models. *Machine Learning*, 37:183–233, 2004.

Kingma, D. P. and Welling, M. Auto-encoding variational bayes. *CoRR*, abs/1312.6114, 2014.

Li, Y. and Mandt, S. Disentangled sequential autoencoder. In *ICML*, 2018.

Liang, K. Y. and Zeger, S. L. Longitudinal data analysis using generalized linear models. *Biometrika*, 73:13–22, 1986.

Locatello, F., Poole, B., Rätsch, G., Scholkopf, B., Bachem, O., and Tschannen, M. Weakly-supervised disentanglement without compromises. *Proceedings of the 37th International Conference on Machine Learning*, 2020.

Németh, J. Adversarial disentanglement with grouped observations. In *AAAI*, 2020.

Pinheiro, J. C. and Bates, D. M. Mixed-effects models in s and s-plus. *Technometrics*, 43:113 – 114, 2001.

Quinn, G. P. and Keough, M. Experimental design and data analysis for biologists. 2002.

Rezende, D. J., Mohamed, S., and Wierstra, D. Stochastic backpropagation and approximate inference in deep generative models. In *ICML*, 2014.

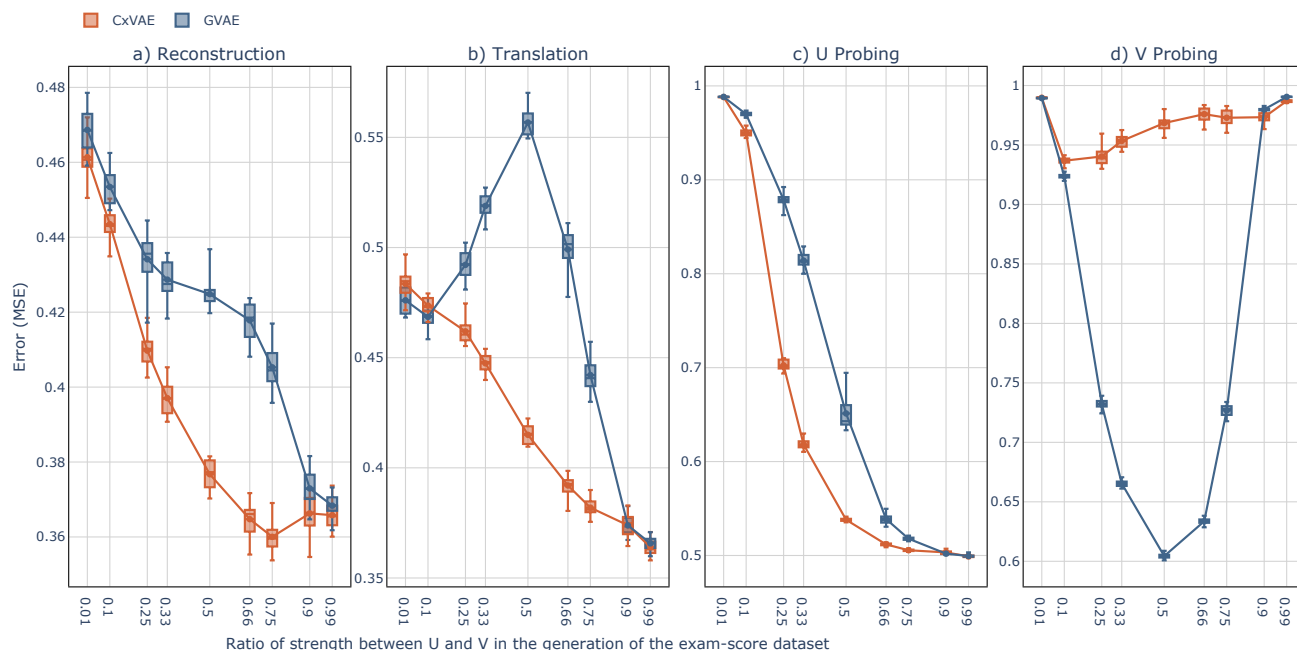


Figure 6. Ablation study. Replacing any novel component of our CxVAE with an existing alternative leads to worse performance. Note that the largest reduction in performance is brought by making the instance encoder unconditional. The next is changing the regularization, followed by using a less expressive group encoder.

Schafer, J. L. and Graham, J. W. Missing data: our view of the state of the art. *Psychological methods*, 7 2:147–77, 2002.

Shu, R., Chen, Y., Kumar, A., Ermon, S., and Poole, B. Weakly supervised disentanglement with guarantees. In *International Conference on Learning Representations*, 2020. URL <https://openreview.net/forum?id=HJgSwyBKvr>.

Tenenbaum, J. B. and Freeman, W. T. Separating style and content with bilinear models. *Neural Computation*, 12: 1247–1283, 2000.

Zaheer, M., Kottur, S., Ravanbakhsh, S., Póczos, B., Salakhutdinov, R., and Smola, A. Deep sets. In *NIPS*, 2017.

A. You *can* have an appendix here.

You can have as much text here as you want. The main body must be at most 8 pages long. For the final version, one more page can be added. If you want, you can use an appendix like this one, even using the one-column format.